

# MDS

June 22, 2025

## Multi-Dimensional Scaling (MDS)

ML2: AI Concepts and Algorithms (SS2025)  
*Faculty of Computer Science and Applied Mathematics*  
*University of Applied Sciences Technikum Wien*



**Lecturer:** Rosana de Oliveira Gomes  
**Authors:** S. Rezagholi, R.O. Gomes



### Presentation – Slide 1: Introduction to Multi-Dimensional Scaling (MDS)

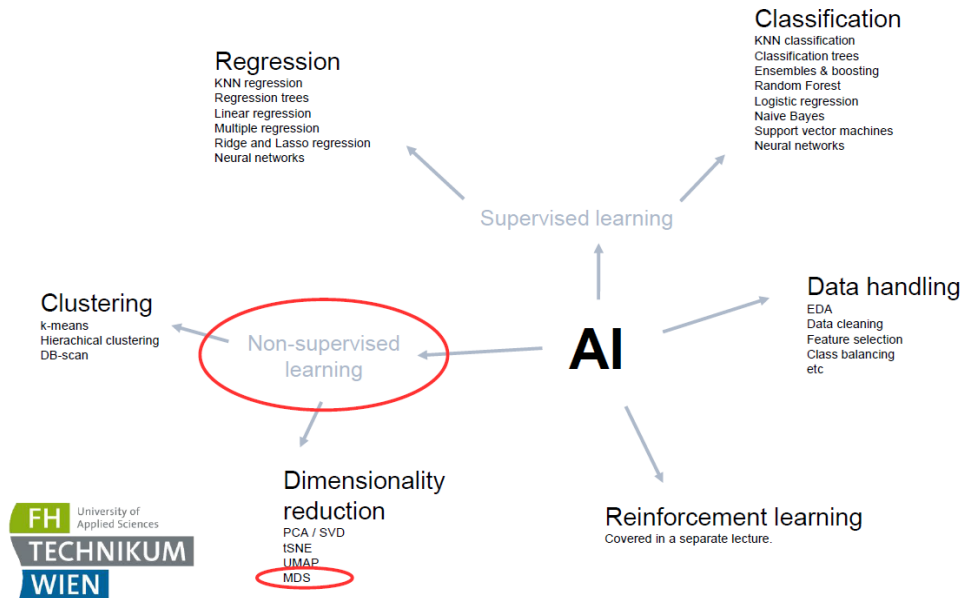
Good [morning/afternoon], everyone.

Welcome to today's presentation on **Multi-Dimensional Scaling (MDS)** — a technique used in machine learning and data analysis to visualize the similarity or dissimilarity between data points in a lower-dimensional space.

This lecture is part of the **ML2: AI Concepts and Algorithms** course for the Summer Semester 2025, presented by the **Faculty of Computer Science and Applied Mathematics** at the **University of Applied Sciences Technikum Wien**.

The material was prepared by **S. Rezagholi and Rosana de Oliveira Gomes**, who is also the lecturer for this session.

Let's now proceed to the next slide — please share it when ready.



## Slide 2: Where Does MDS Fit in AI?

Let's begin by placing **Multidimensional Scaling (MDS)** in the broader context of **Artificial Intelligence**.

AI consists of various learning paradigms, and these are commonly grouped into three main categories:

- **Supervised learning** (e.g., regression and classification),
- **Unsupervised learning** (e.g., clustering and dimensionality reduction),
- **Reinforcement learning** (handled in a separate lecture).

MDS belongs to **unsupervised learning**, as highlighted by the red circle in the bottom-left quadrant.

Under **Dimensionality Reduction**, MDS joins other techniques like:

- **PCA / SVD** (Principal Component Analysis, Singular Value Decomposition),
- **t-SNE**,
- **UMAP**.

The goal of these methods is to **reduce high-dimensional data** into a more **interpretable low-dimensional form**, typically 2D or 3D — while **preserving the structure or relationships** between points as best as possible.

MDS, specifically, focuses on preserving **pairwise distances or dissimilarities** in the transformation process.

When you're ready, please proceed to the next slide.

# Dimensionality Reduction Recap

Dimensionality curse: challenges and complexities that arise when working with high-dimensional data, where the number of features or variables is significantly large

High dimensional systems: e.g. genomics, environmental science, NLP, customer segmentation.

**Dimensionality reduction** transforms data from high-dimensional space into a low-dimensional space while **preserving as much of the dataset's original information as possible**.



*How to convert a multi-dimensional dataset into a small dimension visualization?*




Image Source: [https://www.kdnuggets.com/2019/04/dimensionality-reduction.html](#)

## Common algorithms

- Principle Component Analysis
- T-SNE
- UMAP

## Slide 3: Dimensionality Reduction Recap

Before diving deeper into MDS, let's quickly **recap the concept of dimensionality reduction**.

**The Curse of Dimensionality** refers to the challenges that emerge when working with data that has a **large number of features**. As the number of dimensions increases, data becomes sparse, computations grow more expensive, and patterns become harder to detect.

This issue appears in many **high-dimensional systems**, such as:

- **Genomics**
- **Environmental science**
- **Natural Language Processing (NLP)**
- **Customer segmentation**

**Dimensionality Reduction** techniques aim to **project high-dimensional data into a lower-dimensional space** — typically 2D or 3D — while **preserving the structure, distance, or similarity relationships** as much as possible.

Common algorithms include:

- **PCA** (Principal Component Analysis)
- **t-SNE** (t-distributed Stochastic Neighbor Embedding)
- **UMAP** (Uniform Manifold Approximation and Projection)

The key question these algorithms answer is: **“How can we convert a multi-dimensional dataset into a compact, visual form without losing the essence of the original information?”**

MDS, the focus of this lecture, is another answer to this question — one that prioritizes **pairwise distances or dissimilarities**.

Let's move on when you're ready.

## Recap

**PCA (Principal Component Analysis):** Captures **linear** relationships among variables while maximizing **variance**. It is fast and interpretable, but limited to linear patterns.

**t-SNE (t-Distributed Stochastic Neighbor Embedding):** Focuses on **local** structures and is great for **cluster visualization**. It captures nonlinear relationships, but is computationally intensive and less interpretable.

**UMAP (Uniform Manifold Approximation and Projection):** Balances **local** and **global** structures. It is very fast, flexible, and effective for large, nonlinear data. UMAP maintains more global structure than t-SNE.



### Slide 4: Recap of Common Dimensionality Reduction Techniques

To better understand where **MDS** stands, let's review three widely used dimensionality reduction methods:

---

#### PCA (Principal Component Analysis)

- Captures **linear relationships** between variables.
- Maximizes **variance**, projecting the data into new axes (principal components) that best explain the variability.
- Fast, interpretable, but **limited to linear structures**.

---

#### t-SNE (t-Distributed Stochastic Neighbor Embedding)

- Emphasizes **local structure**, making it ideal for **cluster visualization**.
- Excellent at revealing **nonlinear relationships**, but...
- It is **computationally expensive** and often **less interpretable** due to stochastic behavior and lack of global structure preservation.

---

#### UMAP (Uniform Manifold Approximation and Projection)

- Balances **local and global** structure preservation.
  - More **efficient and scalable** than t-SNE.
  - Captures complex **nonlinear patterns** and tends to preserve the **overall shape** of the data better than t-SNE.
-

Each of these methods has a different philosophy about what should be preserved: **variance**, **local neighborhoods**, or **manifold structure**.

Now, MDS adds another perspective by preserving **pairwise distances** or **dissimilarities** directly. We'll see how it compares next.

Please go ahead with the next slide.


## Multi-Dimensional Scaling

**Goal:** Given pairwise dissimilarities, reconstruct a map that preserves distances.

MDS is a class of algorithms designed to reach an optimal low-dimensional configuration.

Visual representation of distances or dissimilarities between datapoints.

**Use Cases:** relationship among datapoints (e.g. social sciences, epidemiology, etc).



### Slide 5: Multi-Dimensional Scaling (MDS)

Now we arrive at the core topic of today's presentation: **Multi-Dimensional Scaling**, or **MDS**.

---

**Goal of MDS:** Given **pairwise dissimilarities** between data points, MDS aims to reconstruct a **low-dimensional map** such that the distances between points in this new space reflect the **original dissimilarities** as closely as possible.

---

MDS is **not a single algorithm**, but rather a **family of algorithms** designed to find an optimal configuration in 2D or 3D, where the **geometry (distance relationships)** of the data is preserved.

It's often used when we don't have original high-dimensional features — only a **distance or dissimilarity matrix** is available. This makes it especially powerful for **visualizing relationships**.

---

#### Use Cases:

- Social sciences (e.g., opinion similarity)
- Psychology (e.g., perception studies)
- Epidemiology (e.g., genetic or phenotypic similarity)
- Marketing (e.g., brand positioning)

---

There are three main **types of MDS**:

1. **Classical MDS** (also called **Principal Coordinates Analysis**, or **PCoA**):

- Based on **eigen decomposition** of the distance matrix.

2. **Metric MDS**:

- Tries to preserve the actual **distance values**.

3. **Non-metric MDS**:

- Focuses only on **rank order** of dissimilarities, not their exact values.

We'll break these down in upcoming slides.

Let's continue when you're ready.

## Multi-Dimensional Scaling

We are given  $n$  datapoints  $\{x_1, \dots, x_n\}$  with corresponding distance matrix  $D$ .

We want to find a set of  $n$  points  $\{\hat{x}_1, \dots, \hat{x}_n\}$  in  $\mathbb{R}^p$  such that

$$\text{dist}(\hat{x}_i, \hat{x}_j) \approx \text{dist}(x_i, x_j)$$

holds as best as it can.

We want to embed the datapoints into  
*lower-dimensional* space with minimal  
distortion of the distances!



### Slide 6: MDS – The Mathematical Goal

Let's now formalize what **MDS** is trying to achieve.

---

We are given:

- A set of  $n$  data points  $\{x_1, x_2, \dots, x_n\}$
- Along with their **pairwise distances**, stored in a **distance matrix**  $D$

But instead of the original space, we aim to find:

- A new set of points  $\{\hat{x}_1, \hat{x}_2, \dots, \hat{x}_n\}$
  - Embedded in a **lower-dimensional space**  $\mathbb{R}^p$ , where typically  $p = 2$  or  $3$
-

The objective is to ensure that:

$$\text{dist}(\hat{x}_i, \hat{x}_j) \approx \text{dist}(x_i, x_j)$$

That is, the distance between the **mapped points** should approximate the **original distances** as closely as possible.

---

In plain terms:

“We want to embed the datapoints into a **lower-dimensional space**, with **minimal distortion** of the original pairwise distances.”

This is the heart of **MDS** — maintaining the structure of relationships, even after compression into fewer dimensions.

Ready for the next slide!


## Classical Multi-Dimensional Scaling (cMDS)

Given a dataset with respective features.

**Principal Component Analysis** converts correlations into a 2D graph.

**cMDS**: converts **distances** among the samples into a 2D graph. Assumes Euclidean distance matrix.

cMDS is also known as **Principal Coordinate Analysis (PCoA)**



Euclidean distance

$$\sqrt{(3 - 0.25)^2 + (2.9 - 0.8)^2}$$

$$(2.9 - 0.8) = 2.1$$
$$(3 - 0.25) = 2.75$$

### Slide 7: Classical Multi-Dimensional Scaling (cMDS)

Let's now look at the first variant of MDS — **Classical MDS**, also known as **Principal Coordinates Analysis (PCoA)**.

---

#### Contrast with PCA:

- PCA transforms **correlations or covariances** into principal components for dimensionality reduction.
  - In contrast, **cMDS** starts with a **distance matrix**, not feature vectors.
-

## What does cMDS do?

- It takes **pairwise distances** between samples (typically Euclidean)
  - And converts this information into a configuration of points in a lower-dimensional space — usually **2D or 3D**
  - Such that the **geometrical distances between points** approximate the original **distances among the samples**
- 

In the bottom example:

- We have data from two cells across two genes.
- We calculate the **Euclidean distance**:

$$\sqrt{(3 - 0.25)^2 + (2.9 - 0.8)^2} \approx 3.45$$

- This process is repeated for all pairwise combinations, building the full distance matrix.
- 

So essentially:

**Classical MDS = Eigen-decomposition of the distance matrix** Result = a spatial configuration of the points that best preserves those distances.


We'll later see how this differs from **metric** and **non-metric MDS**.

Please continue with the next slide.

## cMDS Example

Minimizing the Euclidean distance is the same as maximizing the correlation. PCA and PCoA are the same if considering linear distances.

*E.g.* MDS creates map of distance across cities.  
Distances relationships are real. Does not know orientations.



### Slide 8: cMDS Example — Map of Cities

Here we have a classic and intuitive example to illustrate how **Classical MDS (cMDS)** works.



---

First, note this theoretical equivalence:

**Minimizing Euclidean distance**   **Maximizing correlation** This is why **PCA** and **PCoA (cMDS)** can produce similar results when we're working with **linear** distances and features.

---

### Example: Mapping Cities

- Suppose we know only the **pairwise distances** between cities in the UK — but we don't know their coordinates.
- cMDS takes this distance matrix and produces a **2D plot** that **preserves those distances** as best as possible.
- The result is a **map-like reconstruction** of the cities' positions.

Look at the plots:

- On the **left**, cMDS output: relative positions between cities based on distances.
  - On the **right**, the actual **geographical map**.
- 

### Important takeaway:

cMDS preserves **distance relationships**, but not **absolute orientation or rotation**. You can **rotate**, **translate**, or **flip** the entire map — the relative distances remain the same.

It's a beautiful demonstration of how **geometry can be recovered from dissimilarities** alone.

Ready for the next slide.

## Classical Multidimensional Scaling

**Algorithm Steps:** from Standardized data

### 1. Distance Matrix

	Standardized data				Distance matrix					
	DBP	SBP	BMI	Chol		1	2	3	4	5
Patient 1	-0.917	-1.086	-0.955	0.000	1	0	1.426	4.356	1.463	2.741
Patient 2	-0.262	-0.926	-0.665	1.222	2	1.426	0	4.327	2.051	2.314
Patient 3	1.703	0.990	1.504	-1.324	3	4.356	4.327	0	3.093	2.866
Patient 4	-0.426	0.032	-0.376	-0.560	4	1.463	2.051	3.093	0	1.809
Patient 5	-0.098	0.990	0.492	0.662	5	2.741	2.314	2.866	1.809	0

$d_{1,3} = \sqrt{((-0.917) - 1.703)^2 + ((-1.086) - 0.990)^2 + ((-0.955) - 1.504)^2 + (0 - (-1.324))^2} = 4.356$

### 2. Double Centering Matrix (B)

Computed from squared distance D.

The centered matrix is computed from D so that the mean of each variable is zero. The double centering matrix is:

$$B = -\frac{1}{2} \mathcal{C} D^{(2)} \mathcal{C}$$

#### Interpretation:

- Transforms distance information into geometry.
- Ensures that the new points are centered at the origin (mean 0)
- Allows to work with distances only to derive coordinates.

$$B = \begin{bmatrix} 2.932 & 1.881 & -4.073 & 0.715 & -1.455 \\ 1.881 & 2.864 & -3.983 & -0.353 & -0.409 \\ -4.073 & -3.983 & 7.897 & -0.517 & 0.676 \\ 0.715 & -0.353 & -0.517 & 0.638 & -0.482 \\ -1.455 & -0.409 & 0.676 & -0.482 & 1.670 \end{bmatrix}$$

## Slide 9: Classical MDS — Step-by-Step Algorithm

Now let's break down the **algorithmic steps** behind **Classical MDS** using standardized data.

---

### 0.0.1 Step 1: Compute the Distance Matrix

- Start with a dataset of standardized features — in this case, values like **DBP**, **SBP**, **BMI**, and **Cholesterol** for 5 patients.
- Then compute the **Euclidean distances** between each pair of data points. Example (highlighted):

$$d_{1,3} = \sqrt{(-0.917 - 1.703)^2 + (-1.086 - 0.990)^2 + \dots} = 4.356$$

- This produces a symmetric **distance matrix**  $D$  with zeros on the diagonal.
- 

### 0.0.2 Step 2: Compute the Double Centering Matrix $B$

- First **square the distance matrix**  $D$ .
- Then apply **double centering** using the formula:

$$B = -\frac{1}{2}CD^{(2)}C$$

where  $C$  is the **centering matrix**, which ensures that the resulting coordinates are **centered at the origin** (mean 0).

---

### 0.0.3 Why this step matters:

- It **transforms distance information into geometry**.
- It **eliminates the need for original coordinates** — distances alone are sufficient.
- The result is a matrix  $B$  that encodes **scalar products** between points in the new space — the foundation for recovering coordinates via **eigendecomposition** (in the next step).

Let's continue to the final steps of the cMDS algorithm when you're ready.

# Classical Multidimensional Scaling

## 3. Calculate Eigenvectors and Eigenvalues

Consider only the first 2 eigenvectors and eigenvalues which will be used for creating the lower dimensional space.

$$E_m = \begin{bmatrix} -0.441 & 0.452 & -0.042 & -0.632 & 0.447 \\ -0.416 & -0.330 & -0.641 & 0.327 & 0.447 \\ 0.783 & 0.219 & -0.359 & -0.498 & 0.447 \\ -0.054 & 0.369 & 0.488 & 0.650 & 0.447 \\ 0.127 & -0.709 & 0.469 & -0.247 & 0.447 \end{bmatrix} \quad \begin{matrix} \lambda_1 = 12.447 \\ \lambda_2 = 2.448 \end{matrix}$$

Represent the eigenvalues in a diagonal matrix:


$$\Lambda_m = \begin{bmatrix} 12.447 & 0 \\ 0 & 2.448 \end{bmatrix}$$



## 4. Calculate Scores

Compute the score matrix that represents the data points in the lower-dimension space:

$$X = E_m \Lambda_m^{1/2} = \begin{bmatrix} -1.556 & 0.706 \\ -1.466 & -0.517 \\ 2.764 & 0.343 \\ -0.190 & 0.577 \\ 0.448 & -1.110 \end{bmatrix}$$



## Slide 10: Classical MDS — Final Steps

We now conclude the **Classical MDS (cMDS)** algorithm with Steps 3 and 4.

### 0.0.4 Step 3: Compute Eigenvectors and Eigenvalues

- Perform **eigendecomposition** of the double-centered matrix  $B$ .
- Select the **top**  $m$  eigenvectors and eigenvalues — typically 2 — to define the new low-dimensional space.
- These represent the **principal axes** of the transformed geometry.

Notation:

- $E_m$ : Matrix of the top  $m$  eigenvectors
- $\Lambda_m$ : Diagonal matrix of the corresponding top  $m$  eigenvalues

### 0.0.5 Step 4: Compute the Coordinates (Scores)

To project the data into the low-dimensional space, we use:

$$X = E_m \Lambda_m^{1/2}$$

- This gives us the **coordinates** of the  $n$  data points in 2D.
- Each row of  $X$  is a data point.
- This is the actual **embedding** we visualize in the MDS plot.

The **MDS plot** (bottom right):

- Plots **MDS1 vs MDS2**, showing the projected data points.
- The axes are scaled by the variance captured (e.g., 78% for MDS1, 15% for MDS2).
- The geometric relationships now reflect the **original distances** between samples.

---

In summary:

Classical MDS turns pairwise distances into coordinates, purely through **distance matrix processing, centering, and eigendecomposition** — no need for the original feature vectors.

Let me know when to proceed with the next slide.

## Metric Multi-Dimensional Scaling

In mathematics, a distance function (that gives a distance between two objects) is also called *metric*, satisfying

- $d(x, y) \geq 0$ , → Distances are positive
- $d(x, y) = 0$  if and only if  $x = y$ , → Distance to itself is zero
- $d(x, y) = d(y, x)$ , → Symmetry
- $d(x, z) \leq d(x, y) + d(y, z)$ . → Triangle Inequality:  
sum of the shorter distances  
between points in a triangle is  
larger than the longest distance.



If the last inequality does not hold, one often speaks of a **dissimilarity** instead of a metric or distance.

### Slide 11: Metric Multi-Dimensional Scaling (MDS)

Now we shift our focus from **classical MDS** to **metric MDS**, which generalizes the idea beyond Euclidean geometry.

---

In mathematics, a **metric** is a function that defines a distance between any two elements in a space — and it must satisfy **four properties**:

#### 1. Non-negativity

$$d(x, y) \geq 0$$

→ Distances are always positive.

#### 2. Identity of indiscernibles

$$d(x, y) = 0 \text{ if and only if } x = y$$

→ An object has zero distance to itself.

### 3. Symmetry

$$d(x, y) = d(y, x)$$

→ The distance is the same in both directions.

### 4. Triangle inequality

$$d(x, z) \leq d(x, y) + d(y, z)$$

→ Any side of a triangle can't be longer than the sum of the other two.

---

If all four of these properties are satisfied, we call it a **metric**. However, if **triangle inequality fails**, we speak of a **dissimilarity** rather than a true distance.

---

In **metric MDS**, we assume that the input dissimilarities behave like **true distances**, even if they're not Euclidean. The goal remains the same:

Find a configuration of points in low-dimensional space whose **distances best match the input dissimilarities**.

This differs from **non-metric MDS**, which we'll explore next — where even the exact distance values aren't trusted, only their **rank order**.

Let me know when you're ready for the next slide.

# Metric Multi-Dimensional Scaling

Metric MDS is a more general method that finds a metric for any distance function instead of assuming an Euclidean distance.

## classical MDS

seeks to find an optimal configuration  $\mathbf{x}_i$  that gives  $d_{ij} \approx \hat{d}_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|_2$  as close as possible.

## Distance Scaling

- Relaxing  $d_{ij} \approx \hat{d}_{ij}$  from cMDS by allowing

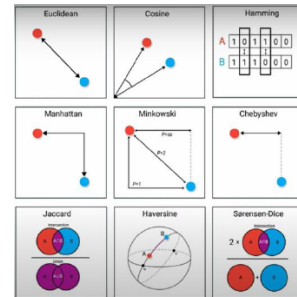
$$\hat{d}_{ij} \approx f(d_{ij}), \text{ for some monotone function } f.$$

## Distance Scaling:

- Metric MDS** if dissimilarities  $d_{ij}$  are **quantitative**.
- Non-metric MDS** if dissimilarities  $d_{ij}$  are **qualitative** (e.g. ordinal).



Unlike cMDS, distance scaling is an optimization process minimizing *Stress function*, and is solved by iterative algorithms.



## Slide 12: Metric Multi-Dimensional Scaling (cont.)

Let's now deepen our understanding of **Metric MDS** by comparing it with Classical MDS and introducing the concept of **distance scaling**.

### 0.0.6 From Classical to Metric MDS

- In **Classical MDS (cMDS)**, we assume **Euclidean distances** and try to satisfy:

$$d_{ij} \approx \hat{d}_{ij} = \|x_i - x_j\|_2$$

- Metric MDS** generalizes this idea:
  - Instead of matching exact Euclidean distances,
  - We allow a **monotonic transformation** of the dissimilarities:

$$\hat{d}_{ij} \approx f(d_{ij})$$

where  $f$  is a monotonic (order-preserving) function.

This flexibility allows Metric MDS to work with **any valid metric** (Euclidean, Manhattan, Cosine, Hamming, etc.).

### 0.0.7 Distance Scaling & Stress

- This process of fitting the distances to the dissimilarities is called **distance scaling**.

- It becomes an **optimization problem**: We try to **minimize a stress function** — a measure of how much distortion exists between input dissimilarities and resulting distances.
- Solved using iterative methods.

### 0.0.8 Key Distinction

- Use **Metric MDS** when dissimilarities  $d_{ij}$  are **quantitative** (e.g., real distances, scores).
- Use **Non-metric MDS** when dissimilarities are **qualitative** (e.g., ranks, preferences).

The image on the right shows various **distance metrics** (Euclidean, Cosine, Hamming, Jaccard, etc.) that can be used in Metric MDS.

Ready for the next slide: **Non-metric MDS**.

## Metric Multi-Dimensional Scaling

**Stress Function: Standardized Residual Sum of Squares**

$$\text{stress} = \mathcal{L}(\hat{d}_{ij}) = \left( \frac{\sum_{i < j} (\hat{d}_{ij} - f(d_{ij}))^2}{\sum_{i < j} d_{ij}^2} \right)^{\frac{1}{2}}$$

Normalization to make the expression scale-invariant.

Stress values are between 0 and 1:

Stress	Goodness of fit
20%	poor
10%	fair
5%	good
2%	excellent
0%	"perfect"

The function  $f$  can be taken to be a parametric monotonic function, such as  $f(d_{ij}) = \alpha + \beta d_{ij}$ .

Metric MDS minimizes the stress as a function of the parameters in  $f$ .



### Slide 13: Metric MDS – The Stress Function

Let's now take a closer look at the **stress function**, which is the core of how Metric MDS works.

### 0.0.9 What is the Stress Function?

The **stress function** quantifies the **discrepancy** between the actual distances in the low-dimensional space ( $\hat{d}_{ij}$ ) and the transformed dissimilarities  $f(d_{ij})$ .

It is defined as:

$$\text{stress} = \mathcal{L}(\hat{d}_{ij}) = \left( \frac{\sum_{i < j} (\hat{d}_{ij} - f(d_{ij}))^2}{\sum d_{ij}^2} \right)^{1/2}$$

- The **numerator** captures how far the embedded distances are from the transformed dissimilarities.
  - The **denominator** is a normalization term — making the expression **scale-invariant** (independent of units or magnitude).
- 

#### 0.0.10 What is $f(d_{ij})$ ?

- $f$  is a **monotonic, parametric function** — often linear:

$$f(d_{ij}) = \alpha + \beta d_{ij}$$

- Metric MDS **optimizes**  $\alpha$  and  $\beta$  (or more complex forms of  $f$ ) to **minimize the stress**.
- 

#### 0.0.11 Interpreting Stress Values

Stress	Fit Quality
> 20%	Poor
~10%	Fair
~5%	Good
< 2.5%	Excellent
0%	Perfect (ideal)

The **lower** the stress, the **better the embedding** preserves the original relationships.

---

So to recap:

**Metric MDS** uses iterative optimization to find a low-dimensional configuration that minimizes **stress** — preserving **quantitative** dissimilarity information as faithfully as possible.

Next up: **Non-metric MDS** — even more flexible. Ready when you are.



# Metric MDS



## Algorithm Steps

1. **Initialization:** initialize points in random positions
2. **Compute Distances:** obtain distance matrix for the configuration
3. **Compute Loss:** Evaluate *Stress function*.
4. **Optimize:** Apply gradient descent to update/ minimize *stress*

Use (a variant of) gradient descent with an appropriate stopping condition.

Note:

- The number of dimensions influences the outcome.
- The parameters of the optimization procedure influence the outcome.
- The random initialization of points influences the outcome.



## Slide 14: Metric MDS – Algorithm Steps

Let's now summarize the **iterative optimization process** that Metric MDS uses to minimize the stress.

### 0.0.12 Algorithm Steps

1. **Initialization** Randomly initialize the positions of the data points in the target low-dimensional space (e.g., 2D or 3D).
2. **Compute Distances** From the current configuration, compute the **Euclidean distances**  $\hat{d}_{ij}$  between all pairs.
3. **Compute Loss (Stress)** Evaluate the **stress function** between the current distances  $\hat{d}_{ij}$  and the transformed dissimilarities  $f(d_{ij})$ .
4. **Optimize** Use **gradient descent** (or a variant) to update the point positions and reduce the stress. Stop when the improvement is below a given threshold or a maximum number of iterations is reached.

### 0.0.13 Important Notes:

- The **number of target dimensions** (e.g., 2, 3) affects the quality of the embedding.
- The **parameters of the optimization** (like learning rate or stopping criteria) influence the result.
- The **initial random layout** of points can impact convergence — multiple runs may yield different local minima.

The plot in the lower right corner shows how **stress decreases** as the **number of dimensions increases**, improving the fit.

The upper right visualizations show how a complex 3D structure is flattened to 2D while preserving neighborhoods — a common MDS application.

Next, we'll look at **Non-metric MDS**, which drops even more assumptions. Ready when you are.

## MDS IMPLEMENTATION

### Parameters

**dissimilarity**: 'euclidean', 'precomputed', etc.

**metric**: Boolean (True = metric MDS, False = non-metric)

**n\_components**: Target dimension (usually 2 or 3)

**max\_iter, eps**: Convergence settings

**random\_state**: For reproducibility

```
from sklearn.manifold import MDS
mds = MDS(n_components=2, metric=True, random_state=42)
X_transformed = mds.fit_transform(distance_matrix)
```



<https://scikit-learn.org/stable/modules/generated/sklearn.manifold.MDS.html>

### Slide 15: MDS Implementation in Python (Scikit-learn)

Let's wrap up by looking at how to **implement MDS in practice** using Python and Scikit-learn.

---

#### 0.0.14 Key Parameters in `sklearn.manifold.MDS`

- **dissimilarity**: Choose how distances are computed. Options:
    - 'euclidean': Compute Euclidean distances internally.
    - 'precomputed': Supply your own distance or dissimilarity matrix.
  - **metric**:
    - True → use **Metric MDS**
    - False → use **Non-Metric MDS**
  - **n\_components**: The number of output dimensions (usually **2** or **3** for visualization).
  - **max\_iter, eps**: Control **convergence criteria** (max iterations, tolerance for stress minimization).
  - **random\_state**: Ensures **reproducibility** of the embedding (important for presentations or comparisons).
-

### 0.0.15 Sample Code

```
from sklearn.manifold import MDS

mds = MDS(n_components=2, metric=True, random_state=42)
X_transformed = mds.fit_transform(distance_matrix)
```

This runs **metric MDS** with 2 output dimensions and a fixed random seed.

---

Full documentation: [Scikit-learn MDS Documentation](#)

---

That concludes the implementation part. Please send the next slide if there's more — or let me know if we're at the end of the presentation.

## Non-Metric MDS

- In some applications (e.g. psychology or marketing) data is often measured on ordinal scales.
- Use case: sound/music/visual perception, facial expression, gene expression, etc
- For ordinal variables there is a version of MDS that is only sensitive to the relative ranks of 'distances'.
- Ordinal MDS is available in Python via: `sklearn.manifold.MDS(metric=False)`.
- Ordinal MDS is an iterative numerical procedure combining steps of gradient descent and isotonic regression.



### Slide 16: Non-Metric MDS

We now come to the most flexible variant: **Non-Metric MDS**, also known as **Ordinal MDS**.

### 0.0.16 When is Non-Metric MDS useful?

In some fields like **psychology**, **marketing**, or **perceptual studies**, we often deal with **ordinal data** — where we know **ranking** (e.g., preferences or similarity order), but not **exact distances**.

### 0.0.17 Use Cases Include:

- Auditory perception (sound/music similarity)

- Visual perception (e.g., facial expressions)
  - Gene expression levels
  - Emotion or preference rating tasks
- 

#### 0.0.18 Key Idea:

- Instead of preserving exact distances, Non-Metric MDS tries to preserve the **order** of dissimilarities:
    - If item A is more similar to B than to C, this should be reflected in the plot.
  - It applies a **monotonic transformation** to the dissimilarities — not necessarily linear — and focuses on **ranking**.
- 

#### 0.0.19 How It Works

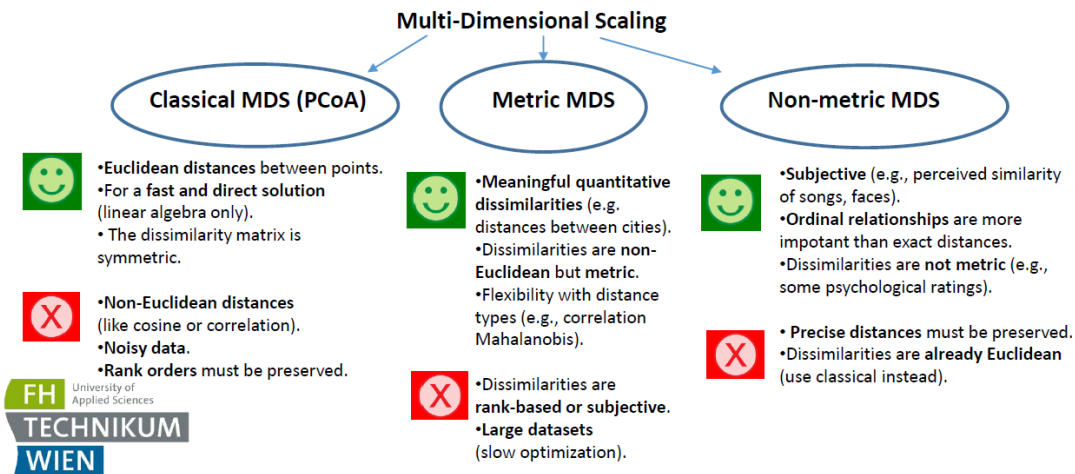
- Available in scikit-learn as:  
`MDS(metric=False)`
  - Internally, it uses:
    - **Gradient descent** to minimize stress
    - **Isotonic regression** to find the best monotonic transformation of dissimilarities
- 

In short:

**Non-Metric MDS** is great when you have only **ordinal data** and want to visualize **relative similarity**, without assuming that your dissimilarities follow any specific metric.

Let me know if there's one final slide (e.g. summary or wrap-up) — or if we should now conclude the presentation.

# When to use what?



## Slide 17: When to Use What?

Let's conclude this presentation with a **comparison guide** to help you choose the right **MDS variant** based on your data and goals.

### 0.0.20 Classical MDS (PCoA)

Use when:

- Dissimilarities are **Euclidean**
- You want a **fast, direct solution** (uses only linear algebra)
- Your distance matrix is **symmetric**

Avoid when:

- Distances are **non-Euclidean** (like cosine)
- The data is **noisy** or contains rank-only information

### 0.0.21 Metric MDS

Use when:

- You have **quantitative**, meaningful dissimilarities (e.g., distances between cities, physical measurements)
- Dissimilarities follow a **metric** but are **not Euclidean**
- You need **flexibility** in distance types (e.g., correlation, Mahalanobis)

Avoid when:

- The dissimilarities are **ordinal** or **subjective**

- Dataset is **very large** (optimization becomes slow)
- 

### 0.0.22 Non-Metric MDS

Use when:

- The dissimilarities are **subjective or ordinal** (e.g., perceived similarity of songs, facial expressions)
- You care about **relative rankings**, not exact distances
- Dissimilarities may **not obey metric properties**

Avoid when:

- You need to **preserve exact distances**
  - Your dissimilarities are **already Euclidean** — then prefer Classical MDS
- 

**Final tip:** Choose your MDS technique based on whether your input dissimilarities are **metric**, **Euclidean**, or **ordinal** — and how important **speed**, **accuracy**, or **interpretability** is in your application.

Thank you for your attention. Let me know if you'd like a summary slide or Q&A support.

## Takeaway

- MDS is a form of nonlinear dimensionality reduction.
- PCA is related to classic MDS, since both use matrix decomposition, but they are not the same: PCA leads to linear dimensionality reduction.
- In certain disciplines it is common to apply (ordinal) MDS to dissimilarities obtained from correlation matrices.
- The concept of Multidimensional Scaling based on elementary mathematics (*Kruskal and Wish 1978*) allow for a more generalized method.

**Next up:** Support Vector Machines



### Final Slide: Takeaway

Let's close with a summary of the key points about **Multidimensional Scaling (MDS)**.

---

### 0.0.23 Key Takeaways

- **MDS** is a form of **nonlinear dimensionality reduction**, especially useful when we want to visualize or interpret **dissimilarities** between objects.
- While **PCA** and **Classical MDS** both rely on **matrix decomposition**, they are **not the same**:
  - **PCA** works on feature covariances → **linear** reduction
  - **MDS** works on distance or dissimilarity matrices → **nonlinear** reduction
- In some fields like **psychology, sociology, or marketing**, it's common to apply **ordinal (non-metric) MDS** to dissimilarities derived from **correlation matrices** or **subjective assessments**.
- The work by **Kruskal and Wish (1978)** laid the foundation for modern MDS — introducing a **generalized, flexible method** for visualizing similarities in data.

---

### Next Topic: Support Vector Machines (SVMs)

We'll now transition to SVMs — a supervised learning technique especially powerful for classification problems in high-dimensional spaces.

Thank you for your attention! Let me know if you'd like a one-slide summary, quiz, or discussion follow-up.